

crsra: A package for Cleaning and Analyzing Coursera Research Export Data

by Aboozar Hadavand, Jeffrey Leek

Abstract An abstract of less than 150 words.

Introduction

It is hard to pin down the time of the birth of the first Massive Open Online Course (MOOC).¹ But since the advent of more focused MOOCs pioneered by universities and platforms such as Coursera, Udacity, and edX, researchers have tried to focus on studying MOOCs. There are fundamental differences between traditional education and MOOCs was large enough to attract researchers to study students' behavior and outcomes. These differences are best reflected in the definition of MOOCs by McAuley et al. (2010) that "[a]n online course with the option of free and open registration, a publicly shared curriculum, and open-ended outcomes which integrates social networking, accessible online resources ... and most significantly builds on the engagement of learners who self-organize their participation according to learning goals, prior knowledge and skills, and common interests."

Research on MOOCs few years with more data being accumulated and collected. Bozkurt et al. (2017) studied literature published on MOOCs through 2015 and found that the number of articles published on the subject increased from 1 in 2008 to 170 in 2015. More research is needed to fully understand the effectiveness, reach, limits, and the potential of MOOCs. However, one of the main challenges in studying MOOCs remains to be data. Data is not usually publicly available since it is owned by private MOOC providers and there are concerns about privacy of students. More importantly, as Lopez et al. (2017) point out, the size and complexity of MOOC data is an overwhelming challenge to many researchers. Therefore, it is imperative to provide tools that pave the way for more research on the new subject of MOOCs.

This paper introduces a package called *crsra* based on the statistical software R to help clean and analyze large loads of data from the Coursera MOOCs. The advantages of the package are as follows: a) faster loading of data for analysis, b) efficient method for combining data from multiple courses and even across institutions,² and c) provision of a set of functions for analysing student behaviors.

Coursera Research Data

Coursera is one of the main providers of MOOCs that launched in January 2012. In fact, with over 25 million learners, Coursera is the biggest provider in the world being followed by EdX, the MOOC provider that was a result of a collaboration between Harvard University and MIT, with over 10 million users. Coursera has over 150 university partners from 29 countries and offers a total of 2000+ courses from computer science to philosophy (cou). In addition, Coursera offers 180+ specialization, Coursera's own credential system, and 4 fully online Masters degrees. Courses include recorded video lectures, graded assignment, quizzes, and discussion forums.

Since the early years of the platform, Coursera has encouraged researchers to analyze students' data and has facilitated the use of the data and the platform for A/B testing. Starting November 2015 Coursera introduced a dashboard for self-service data exports. Through this tool, partner institutions and instructors can download data for a single course or all courses associated with the institution. Research data exports are sets of CSV files and are designed for use in relational database systems. One of the advantages of the data is the existence of a single *hashed user ID* for each student. This user ID is consistent for learners across all courses offered by an individual institution and allows for connecting learner grades and progress across course.

There are five types of research data export for each course. The Table 1 summarizes these five types. This set of data is written in roughly 100 tables: some containing course information and content, some containing students' information, progress, and outcomes, and some containing forum data. Figure 1 shows

¹Some have claimed Sesame Street as the first MOOC. Delaney Parrish, "Sesame Street was the original MOOC," *BROOKINGS NOW*, The Brookings Institution, June 18, 2015, <https://www.brookings.edu/blog/brookings-now/2015/06/18/sesame-street-was-the-original-mooc/>

²This is important since although MOOC researchers have access to thousands of students in their sample, few studies benefit from data across multiple courses and institutions. Such analysis helps draw more robust conclusions about student behaviors (Reich, 2015).

Table 1: Types of research data export

Data Type	Description
Assessment submission data	Assessment submissions of quizzes, peer review, and programming assignments by learners.
Course grade data	Contains the highest grade achieved by each learner on each required assessment as well as the timestamp of the learner’s highest-scoring submission. This table also includes each learner’s overall grade in the course.
Course progress data	Contains data data documenting the timestamp for when the learner interacted with each piece of course content and the timestamps for when items were opened, completed, reopened, reattempted, etc.
Demographic data	Contains the following information for all enrolled learners: general geographical data (based on IP address), browser language preference, and information for learners who completed their learner profile responses or participated in Coursera’s platform-wide demographic survey (including age, gender, education level, and employment status).
Discussion data	Contains forum activity data such as posts, responses, upvotes/downvotes, flags, and questions and answers associated with course content items.

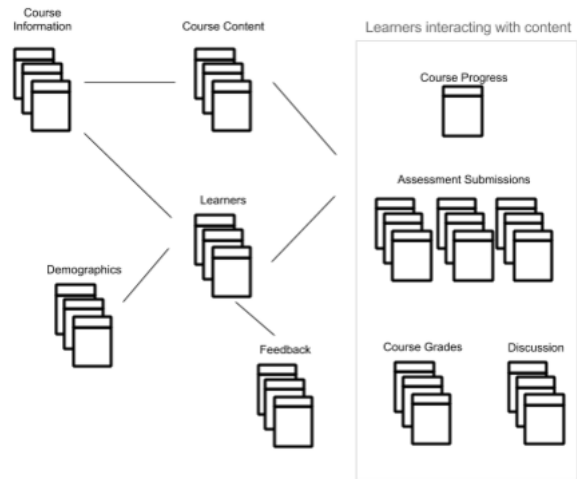


Figure 1: The major relationships between tables groups, with minor connections omitted (Source: Coursera)

While Coursera provides tools for creating Postgres databases in a docker container³, as we mentioned earlier, importing data for analysis remains to be a challenge for researchers with limited experience with relational databases. Moreover, such tools are usually not platform independent.⁴

The *crsra* Package

The *crsra* package helps import and organize Coursera's research data exports into R. It also run some preliminary analysis on the data. In the following section, we introduce the package and provide instruction on how to import Coursera research data exports. To install this package, you will need to install *devtools*. Install the [devtools package](#), available from CRAN. Then execute the following commands to install the *crsra* package

```
library("devtools")
devtools::install_github("jhudsl/crsra", build_vignettes = TRUE)
```

In order to import your data dump into R, first point your working directory to the directory that contains all the unzipped course folders. Then execute the command `crsra_import()`. If you are not pointing to the correct directory, you will receive a warning and the execution will be halted. Note that the data import may take some time if the course data is large and there are several courses in your working directory. Also note that by running the `crsra_import()` command, you import all tables for each individual course into R in a list called `all_tables`.

For a list of all the tables in the data download, please click [here](#). All tables can be called using `all_tables[["course_name"]][["table_name"]]`. For instance, if you like to call the table `peer_comments` in the course `Regression Models`, you can simply execute `all_tables[["Regression Models"]][["peer_comments"]]`. To see a list of courses imported by the `crsra_import()` command check the variable `coursenames`. To see a list of all the tables check the variable `tablenames`.

To see the data import in use, we use the package on data from Johns Hopkins University (JHU) Data Science Specialization on Coursera. This specialization, developed by Jeffrey Leek, Roger Peng, and Brian Caffo, consists of ten courses. There has been more than two million enrollments since the launch of this program in April 2014. The size of data on the students who took these ten courses since 2015 is around 18 gigabytes. We used the *crsra* package to import the data on all the courses and then to find the number of students who passed a specific course item (course item 67c10) in the course "Regression Models" and their average grade in a specific course.

```
library(dplyr)

all_tables[["Regression Models"]][["course_item_grades"]] %>%
  dplyr::filter(course_item_id == "67c10") %>%
  dplyr::filter(course_item_passing_state_id == 2) %>%
  dplyr::summarise(n = n(), grade = mean(course_item_grade_verified))

## A tibble: 1 x 2
##       n      grade
##   <int>   <dbl>
## 1  8640 0.9556052
```

The package also includes a few other functions are added to the package in addition to the main `crsra_import()` function. A list of functions and their descriptions is provided in Table 2.

use another example given the functions above

A Preliminary Analysis of Student Behavior on Coursera

Understanding how students progress through an education program is critical for any educational planning and decision making (King, 1972). Models of student progress are needed in order to estimate the probability of a student completing a certain item in a course and predict the time required to finish a course. Furthermore, common measures of academic success and progress cannot be defined in the same way for MOOCs. For instance, as Perna et al. (2014) states, we have limited knowledge on whether learners' progress through a MOOC should be measured in a sequential fashion or in a way that captures the flexibility and freedom in learning behavior that is unique to MOOCs.

³The tools is called 'courseraresearchexports' and can be found here: <https://github.com/coursera/courseraresearchexports>

⁴In an initial version of *crsra* based on Postgresql we had the problem of some team members not being able to set up the database properly on their PCs.

Table 2: Other functions in the **crsra** package

Function	Description
<code>crsra_membershares</code>	Returns a summary of the total number and the shares of users in each course broken down by factors such as roles, country, language, gender, employment status, education level, and student status.
<code>crsra_gradesummary</code>	Returns total grade summary or broken down by the factors mentioned above.
<code>crsra_progress</code>	Summarizes, for each course item, the total number and the share of users who stopped the course at that specific course item. The function ranks course items by their attrition.
<code>crsra_assessmentskips</code>	Users may "skip" reviewing a submission if there is a problem with it. This function categorizes skips by their type such as "inappropriate content", "plagiarism", etc. The function also returns list of mostly used words in peer comments.
<code>crsra_timetofinish</code>	Calculates the time to finish a course for each user.

We have limited understanding of user progress in MOOCs. There are only a handful of studies on the subject of student pace, who completes classes, and learning sequence in MOOCs. [Perna et al. \(2014\)](#) perform a descriptive analysis of student progress through a set of 16 courses on Coursera. They find that most users accessed course content in the sequential order defined by the instructor of the course. [Ho et al. \(2014\)](#) study 17 courses taught on EdX and find that most of the attrition in courses happen in the first week of course activity (about 50 percent attrition) and that the average percentage of learners who cease activity in the second week declines sharply to 16 percent. Most of these studies are specific to a set of courses or platforms. Due to the many differences in the characteristics of MOOCs, any extrapolation of the results to MOOCs in general has to be done with caution.

In the following section, we will investigate students' progress through the ten Data Science Specialization courses on Coursera provided by JHU. Using the `crsra_timetofinish` function in the *crsra* package, we can first investigate the time difference between the first and last activities within a course for each student. Time to finish is only calculated for those who finished the course. Figure 2 depicts the density function for time to finish calculated for three of the courses in the specialization. Note that the density functions vary across courses. While for *Developing Data Products* and *Getting and Cleaning Data* a majority of students finish the courses in around 30 days, for *Data Science Capstone* a majority of students finish the course in 50 days.

```
TTF <- crsra_timetofinish()
```

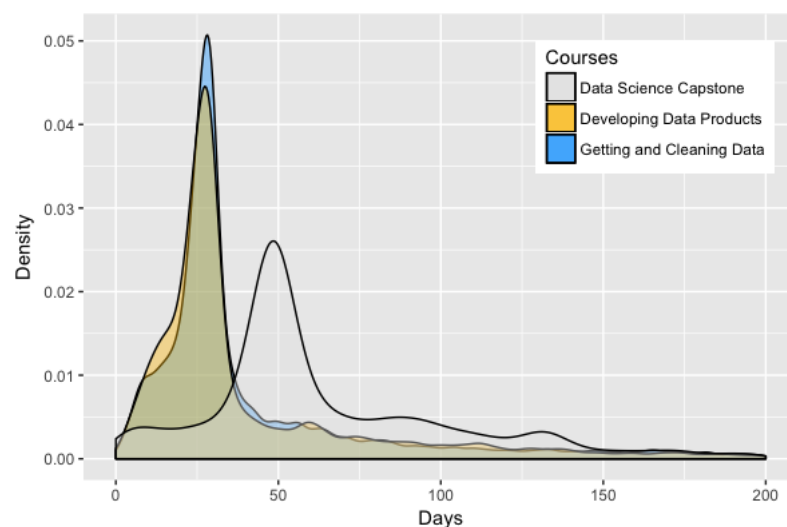


Figure 2: Density functions for time to finish defined as the time difference between the first and last activities across three courses

In the table called `users`, Coursera provides a field for student status of the learner including full-time and par-time students and those who are not degree students. We can look at how time to finish is different for groups with different student status. Figure 3 reports this for the course *Getting and Cleaning Data* and shows that part-time students take longer to finish the course.

```
TTF.Status <- TTF[["Getting and Cleaning Data"]] %>%
  dplyr::left_join(all_tables[["Getting and Cleaning Data"]][["users"]],
    by = "jhu_user_id", `copy`=TRUE) %>%
  dplyr::filter(!is.na(student_status))
```




Figure 3: Density functions for time to finish across groups with different student statuses through the course *Getting and Cleaning Data*

Our next step is to understand student progress in the courses. One of the factors that distinguishes MOOCs from traditional classrooms is the flexibility in advancing through the course. While in traditional education the class length, pace, and completion dates are determined by instructor of the course, in MOOCs it is the student who, for the most part, has the freedom to choose these factors. This is evident by looking at the time gap between completing each item in the course. We can first look at how many course items students pass in the first week of course activity. One obvious but yet interesting finding is large variations across students. For instance, if we look at the course *Getting and Cleaning Data*, we can use the following code to find the number of course items completed in the first week of course activity. The course has roughly 40 items including lectures and assignments. The variable `nweek1` captures the number of passed course items in the first week of course activity, calculated as one week after the first activity in the class. The density function in Figure 4 represents the variations across students. For a majority of learners, the number of passed course items in the first week is two. However, the number of those who finish more than 10 items in the first course is significant. Also interesting is the double-peak shape of the density function. It is interesting to see that there are more people who complete 12 course items in the first week than there are who complete 7 course items. This indicates an interesting structural change in students' pace between course items 7 and 12.

```
passed.items <- all_tables[["Getting and Cleaning Data"]][["course_progress"]] %>%
  dplyr::group_by(jhu_user_id) %>%
  # 604800 is the number of seconds in a week
  dplyr::filter(course_progress_ts <= min(course_progress_ts) + 604800) %>%
  dplyr::summarise(nweek1 = n())
```

A third interesting variable to look at when studying students' progress in MOOCs is the time gaps between each session. In this exercise, we looked at the time lapsed between each two consecutive course items for each learner throughout the course *Getting and Cleaning Data*. We used the following code for the analysis. Note that we have ranked course items based on the timestamp when student passes them and not their natural order defined by the course instructor.

```
gaps <- all_tables[["Getting and Cleaning Data"]][["course_progress"]] %>%
  # 2 is an indicator that the course item is completed
```




Figure 4: Density functions for the number of passed items in the first week of course activity for Getting and Cleaning Data

```
dplyr::filter(course_progress_state_type_id == 2) %>%
dplyr::group_by(jhu_user_id, course_item_id) %>%
# This is for keeping only the latest event for each course item
dplyr::filter(course_progress_ts == max(course_progress_ts)) %>%
dplyr::ungroup() %>%
dplyr::arrange(jhu_user_id, course_progress_ts) %>%
dplyr::group_by(jhu_user_id) %>%
# This is for converting the time gap to hours
dplyr::mutate(time.dif = as.numeric(course_progress_ts -
                                     lag(course_progress_ts))/3600) %>%

dplyr::filter(!is.na(time.dif)) %>%
dplyr::filter(time.dif != Inf | time.dif != -Inf)
```

Figure 5 shows student progress in the course Getting and Cleaning Data for three sample students. The vertical axis is the gap between two consecutive sessions in hours. These three students are chosen intentionally to show three different learning paths. Panel A shows progress for a student with short gaps between sessions for the first half of the course and longer gaps towards the end. For future reference we call this pattern “slowing down” pattern. This pattern is typical of many students in this course. Panel B shows progress for a student with short gaps between sessions in the beginning and the end of the course and longer gaps in the middle. Students in this group are not as common as the first group. Finally, Panel C shows progress for a student with no clear pattern in their progress throughout the course. Only a small group of students follow this pattern in our data.




Figure 5: Student progress through the course Getting and Cleaning Data for three sample students. The vertical axis shows the time gap between completing an item and the next item in hours.

For the average student, we can look at is the average gap between sessions for the first and the second half of the course. We can then calculate how much the average session gap changes from the first half to the second. Across our sample of students who registered for the course, the average change in session gap from the second half to the second half is positive and equal to 132 percent. In other words the gap between session more than doubles from the first half of the course to the second half. Figure 6 shows the density function of this statistic across our sample. The long right tail in the Figure supports the fact that most students follow the slowing-down pattern. However, the Figure also shows that there are students who speed up during the second half of the semester.




Figure 6: Percentage change in time gaps in course progress between the first and the second half

We can do this analysis for some subgroups of our sample. Some of the most interesting categories are gender, educational attainment, and whether the learner paid for taking the course. The variable `was_payment` in the table `users_courses__certificate_payments` captures whether the learner has ever paid for the eligibility of a course certificate, enrolled in the course, and has not made a refund. This purchase could be a “single payment” for the course or a “bulk payment” for a specialization that contains the course. The following code is used for the analysis of payers and non-payers and the results for all categories are shown in Table 3.

```
gaps.payment <- gaps %>%
  dplyr::group_by(jhu_user_id) %>%
  dplyr::summarise(avgtime = mean(time.dif)) %>%
  dplyr::inner_join(all_tables[["Getting and Cleaning Data"]][["course_grades"]],
    by = "jhu_user_id", `copy`=TRUE) %>%
  dplyr::filter(course_passing_state_id %in% c(1, 2)) %>%
  dplyr::left_join(all_tables[["Getting and Cleaning Data"]][["users_courses__certificate_payments"]],
    by = "jhu_user_id", `copy`=TRUE) %>%
  dplyr::filter(!is.na(was_payment)) %>%
  dplyr::group_by(was_payment) %>%
  dplyr::summarise(avggap=mean(avgtime))
```

Last exercise in our analysis is to study how Coursera’s change in policy from a pay-per-course business model to a subscription model changed students’ progression throughout the course. In October 30, 2016, Coursera introduced a new payment system through which they allowed students to purchase access to all content in a specialization on a month-by-month or annual basis. As a result, student would only pay for the amount of time they need to learn the material. This system replaced the existing model where students would pay up front for each course regardless of how long it took them to finish the course. An interesting exercise then is whether the switch to this system where payments is tied to the length of time it takes students to finish the class make students finish faster. The following code calculates the average number of courses passed in the first week of activity for the two groups: those who enrolled in the course before October 30, 2016 and those who enrolled after. Our hypothesis is that those who pay monthly are more likely to finish more items in the first week than those who pay a fixed price.

```
passed.items.policy <- passed.items %>%
  dplyr::left_join(all_tables[["Getting and Cleaning Data"]][["course_memberships"]],
```

Table 3: Gaps between sessions for different subgroups of learners in Getting and Cleaning Data Course

Categories	Average gap in hours
Gender	Female: 39 Male: 36
Educational Attainment	Less than high school: 152 High school diploma: 39 College (no degree): 39 Associate degree: 23 Bachelor's degree: 32 Master's degree: 38 Professional degree: 22 Doctoral degree: 34
Paid for the course?	Yes: 36 No: 39


```

    by = "jhu_user_id", `copy`=TRUE) %>%
dplyr::filter(!is.na(course_membership_ts)) %>%
dplyr::mutate(subscription = ifelse(course_membership_ts < "2016-11-01 00:00:00",
                                   "before", "after")) %>%

dplyr::group_by(subscription) %>%
dplyr::summarise(subnw = mean(nweek1))

```

The results suggest that those who enrolled before the policy change on average passed three courses less than the group who enrolled after (9 versus 12). Figure 7 shows the density function of the number of passed items in the first week of activity for the two groups. This comparison, however, has a caveat: there is some selection bias since those who enrolled before and those who enrolled after October 2016 may be fundamentally different.

**Figure 7:** Density functions for the number of passed items in the first week of course activity for Getting and Cleaning Data for those who enrolled in the course before October 30, 2016 and those who enrolled after

Discussion

Bibliography

Coursera, about page. <https://blog.coursera.org/about/>. Accessed: 2018-02-02. [p1]

A. Bozkurt, E. Akgün-Özbek, and O. Zawacki-Richter. Trends and patterns in massive open online

- courses: Review and content analysis of research on moocs (2008-2015). *The International Review of Research in Open and Distributed Learning*, 18(5), 2017. [p1]
- A. D. Ho, J. Reich, S. O. Nesterko, D. T. Seaton, T. Mullaney, J. Waldo, and I. Chuang. Harvardx and mitx: The first year of open online courses, fall 2012-summer 2013. 2014. [p4]
- M. A. King. *Primary and secondary indicators of education*. Department of Applied Economics, University of Cambridge, 1972. [p3]
- G. Lopez, D. T. Seaton, A. Ang, D. Tingley, and I. Chuang. Google bigquery for education: Framework for parsing and analyzing edx mooc data. In *Proceedings of the Fourth (2017) ACM Conference on Learning@ Scale*, pages 181–184. ACM, 2017. [p1]
- A. McAuley, B. Stewart, G. Siemens, and D. Cormier. The mooc model for digital practice. 2010. [p1]
- L. W. Perna, A. Ruby, R. F. Boruch, N. Wang, J. Scull, S. Ahmad, and C. Evans. Moving through moocs: Understanding the progression of users in massive open online courses. *Educational Researcher*, 43(9):421–432, 2014. [p3, 4]
- J. Reich. Rebooting mooc research. *Science*, 347(6217):34–35, 2015. [p1]

Aboozar Hadavand
Bloomberg School of Public Health, Johns Hopkins University
615 N. Wolfe Street
Baltimore, MD 21205, USA
hadavand@jhu.edu

Jeffrey Leek
Bloomberg School of Public Health, Johns Hopkins University
615 N. Wolfe Street
Baltimore, MD 21205, USA
jtleek@jhu.edu